

Group_Project

Kyle Ayiku

2023-03-25

Regression

This is the data set that I chose for this project.

```
library(ggplot2)

## Warning: package 'ggplot2' was built under R version 4.2.3

# Reading in the data set.
setwd("C:/Users/amark/Downloads")
housing <- read.csv("housing.csv")
```

Dividing the data into train and test (Part A)

```
# Eliminating the NAs from the dataset.
housing <- housing[complete.cases(housing),]

set.seed(1234)
num <- sample(1:nrow(housing), 0.80 * nrow(housing), replace = FALSE)
train <- housing[num,]
test <- housing[-num,]

sapply(train, function(x) sum(is.na(x) == TRUE))

##          longitude          latitude housing_median_age
total_rooms
##              0              0              0
0
##    total_bedrooms    population      households
median_income
##              0              0              0
0
## median_house_value    ocean_proximity
##              0              0
```

I split the data into 80% train and 20% test and also checked to see if there any null values and unfortunately, I can see that total_bedrooms has 165 missing values.

```
library(dplyr)

## Warning: package 'dplyr' was built under R version 4.2.3
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(janitor)

## Warning: package 'janitor' was built under R version 4.2.3

##
## Attaching package: 'janitor'

## The following objects are masked from 'package:stats':
##
##   chisq.test, fisher.test

# Cleaning the training data.
clean <- clean_names(train)
colnames(train)

## [1] "longitude"      "latitude"        "housing_median_age"
## [4] "total_rooms"    "total_bedrooms"  "population"
## [7] "households"     "median_income"   "median_house_value"
## [10] "ocean_proximity"
```

I attempted to see if there is any of the training data that needs to be cleaned.

Exploratory Data Analysis (Part B)

Brief look at the data.

head(train) # Displays the first few rows of the training data.

```
##      longitude latitude housing_median_age total_rooms total_bedrooms
## 7534   -118.22    33.91             31         571         153
## 8103   -118.22    33.80             36        1285         347
## 7242   -118.12    34.00             31        3281         768
## 8173   -118.12    33.80             36        1257         205
## 9288   -122.55    38.07             38        3392         709
## 627    -122.18    37.70             36        2639         533
##      population households median_income median_house_value
##      ocean_proximity
## 7534           841         158         2.6154         89200         <1H
## OCEAN
## 8103          1291         337         3.7708        157100         NEAR
## OCEAN
## 7242          2385         733         2.7308        173800         <1H
## OCEAN
```

```
## 8173      530      211      5.3701      251400      <1H
OCEAN
## 9288      1894      713      3.0573      350800      NEAR
BAY
## 627       1209      519      4.0268      205500      NEAR
BAY
```

```
dim(train) # How many rows and columns are in the training data.
```

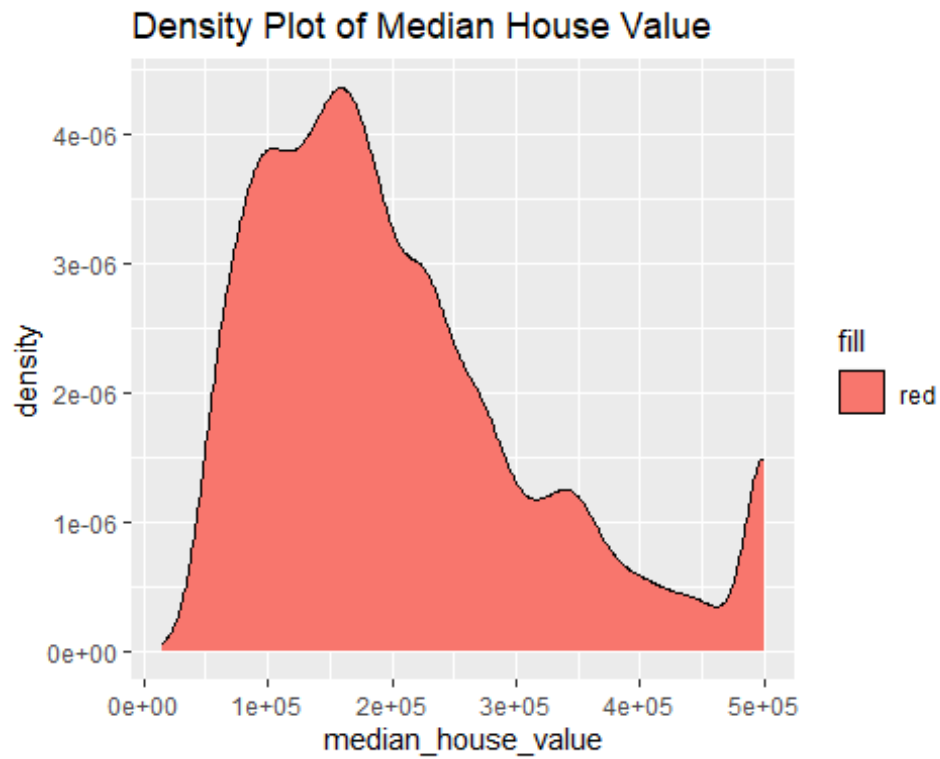
```
## [1] 16346    10
```

```
summary(train) # General statistics of the training data.
```

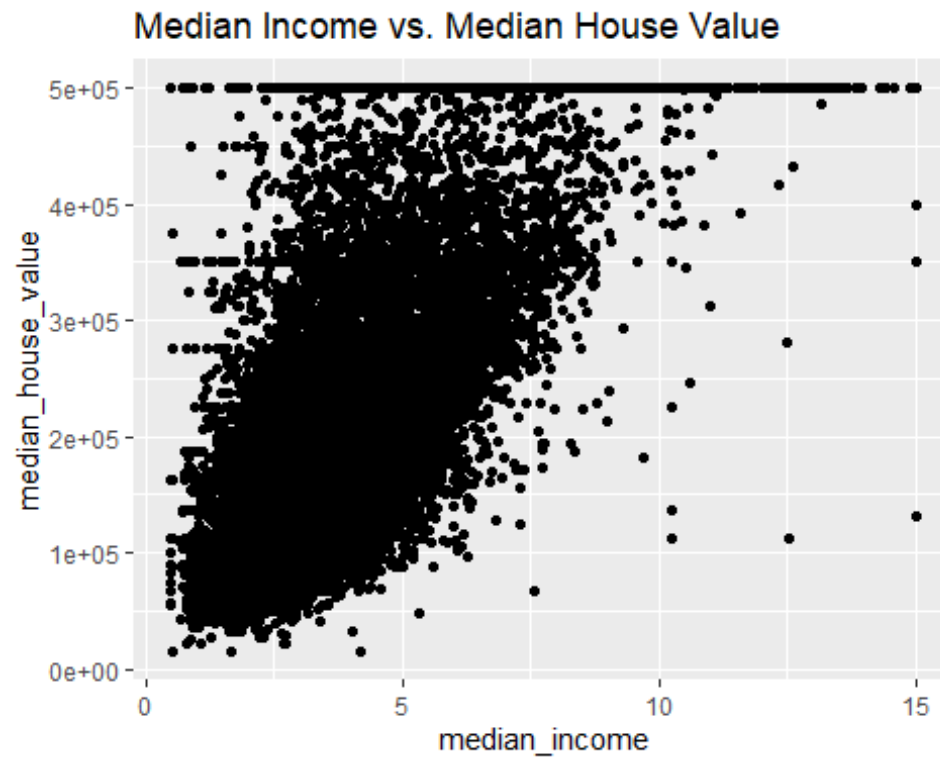
```
##      longitude      latitude      housing_median_age      total_rooms
## Min.   :-124.3    Min.    :32.54    Min.     : 1.00     Min.      :    2
## 1st Qu.: -121.8    1st Qu.:33.93    1st Qu.:18.00     1st Qu.: 1451
## Median : -118.5    Median :34.25    Median :29.00     Median : 2126
## Mean   : -119.6    Mean     :35.62    Mean     :28.59     Mean     : 2637
## 3rd Qu.: -118.0    3rd Qu.:37.71    3rd Qu.:37.00     3rd Qu.: 3138
## Max.   : -114.3    Max.     :41.95    Max.     :52.00     Max.     :39320
## total_bedrooms      population      households      median_income
## Min.    :    1.0    Min.     :    5    Min.     :    1.0    Min.      : 0.4999
## 1st Qu.: 297.0     1st Qu.:  788    1st Qu.: 281.0     1st Qu.: 2.5550
## Median : 435.0     Median : 1166    Median : 409.0     Median : 3.5313
## Mean    : 538.4     Mean      :1425    Mean      :499.5     Mean      : 3.8674
## 3rd Qu.: 647.0     3rd Qu.: 1718    3rd Qu.: 604.0     3rd Qu.: 4.7437
## Max.    :6445.0     Max.      :35682   Max.      :6082.0    Max.      :15.0001
## median_house_value  ocean_proximity
## Min.      :14999      Length:16346
## 1st Qu.:119400      Class :character
## Median :179650      Mode  :character
## Mean      :206728
## 3rd Qu.:264575
## Max.      :500001
```

```
# Density plot for the response variable, "median_house_value."
```

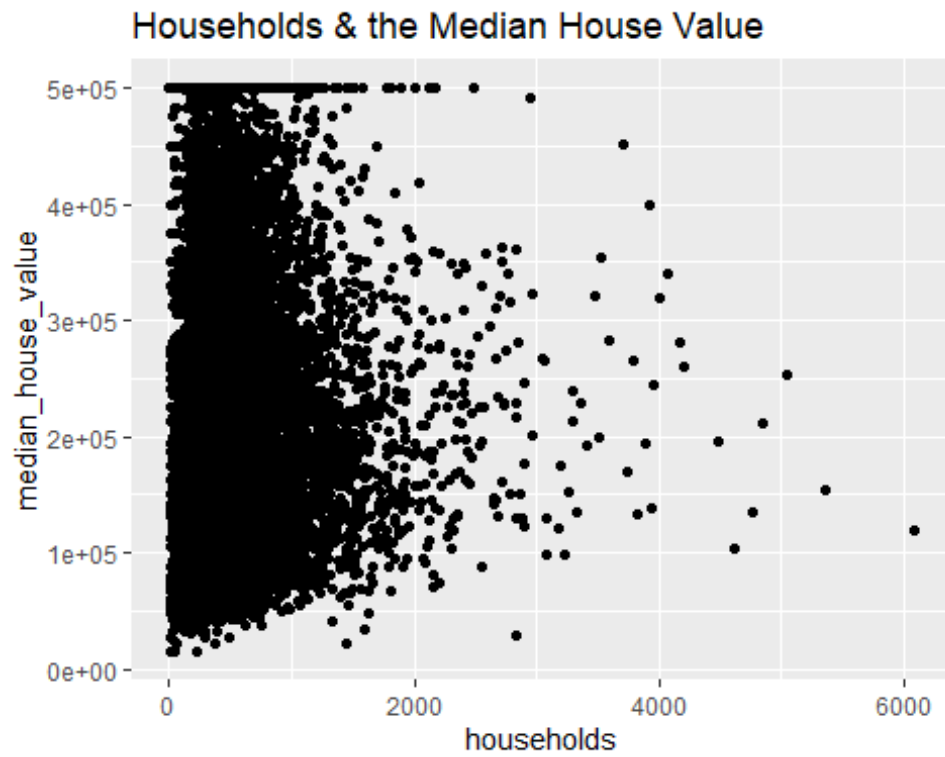
```
ggplot(data = train) +
  geom_density(mapping = aes(x = median_house_value,
                             fill = "red")) +
  labs(title = "Density Plot of Median House Value")
```




```
# Scatter plots for some predictors and median_house_value.  
ggplot(data = train) +  
  geom_point(mapping = aes(x = median_income,  
                           y = median_house_value)) +  
  labs(title = "Median Income vs. Median House Value")
```



```
ggplot(data = train) +  
  geom_point(mapping = aes(x = households,  
                           y = median_house_value)) +  
  labs(title = "Households & the Median House Value")
```



```
ggplot(data = train) +  
  geom_point(mapping = aes(x = population,  
                           y = median_house_value)) +  
  labs(title = "Population vs. Median House Value")
```



Displaying the correlation matrix.

```
library(Hmisc)
```

```
## Warning: package 'Hmisc' was built under R version 4.2.3
```

```
##
```

```
## Attaching package: 'Hmisc'
```

```
## The following objects are masked from 'package:dplyr':
```

```
##
```

```
##      src, summarize
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      format.pval, units
```

```
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 4.2.2
```

```
## corrplot 0.92 loaded
```

```
cor(train[,c(3:9)])
```


```
##              housing_median_age total_rooms total_bedrooms
population
## housing_median_age      1.0000000 -0.3595297 -0.318486183 -
0.291600322
## total_rooms            -0.3595297  1.0000000  0.930276347
```

```

0.855524116
## total_bedrooms          -0.3184862    0.9302763    1.000000000
0.874904940
## population              -0.2916003    0.8555241    0.874904940
1.000000000
## households              -0.3000743    0.9178393    0.978089318
0.906235866
## median_income           -0.1176870    0.1982394   -0.008214284
0.005999249
## median_house_value      0.1071152    0.1354307    0.051267879 -
0.023011587
##
## households median_income median_house_value
## housing_median_age -0.30007428 -0.117687025    0.10711521
## total_rooms        0.91783925  0.198239385    0.13543067
## total_bedrooms     0.97808932 -0.008214284    0.05126788
## population         0.90623587  0.005999249   -0.02301159
## households         1.00000000  0.013615731    0.06699831
## median_income      0.01361573  1.000000000    0.68588807
## median_house_value 0.06699831  0.685888071    1.00000000

corrplot(cor(train[,c(3:9)]), method = "square")

```



Regression Techniques (Part C)

Linear Regression


```

# Creating a multiple linear regression model with
# the target, median_house_value.

lm1 <- lm(median_house_value ~ housing_median_age + total_rooms +
          total_bedrooms + population + households +
          median_income + longitude + latitude, data = train)
summary(lm1)

##
## Call:
## lm(formula = median_house_value ~ housing_median_age + total_rooms +
##     total_bedrooms + population + households + median_income +
##     longitude + latitude, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -559038  -43890  -11392   30223   796426
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -3.575e+06  7.088e+04  -50.438  < 2e-16 ***
## housing_median_age  1.172e+03  4.880e+01   24.019  < 2e-16 ***
## total_rooms      -7.925e+00  8.942e-01   -8.863  < 2e-16 ***
## total_bedrooms    1.082e+02  7.571e+00   14.294  < 2e-16 ***
## population      -3.851e+01  1.210e+00  -31.836  < 2e-16 ***
## households        5.362e+01  8.220e+00    6.523  7.08e-11 ***
## median_income     3.993e+04  3.772e+02  105.874  < 2e-16 ***
## longitude       -4.267e+04  8.081e+02  -52.802  < 2e-16 ***
## latitude        -4.258e+04  7.634e+02  -55.780  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 70010 on 16337 degrees of freedom
## Multiple R-squared:  0.6331, Adjusted R-squared:  0.6329
## F-statistic: 3523 on 8 and 16337 DF, p-value: < 2.2e-16

# Test data evaluation.
pred1 <- predict(lm1, newdata = test)
print(cor1 <- cor(pred1, test$median_house_value)) # Correlation.

## [1] 0.8077368

print(mse1 <- mean((pred1 - test$median_house_value)^2)) # The MSE.

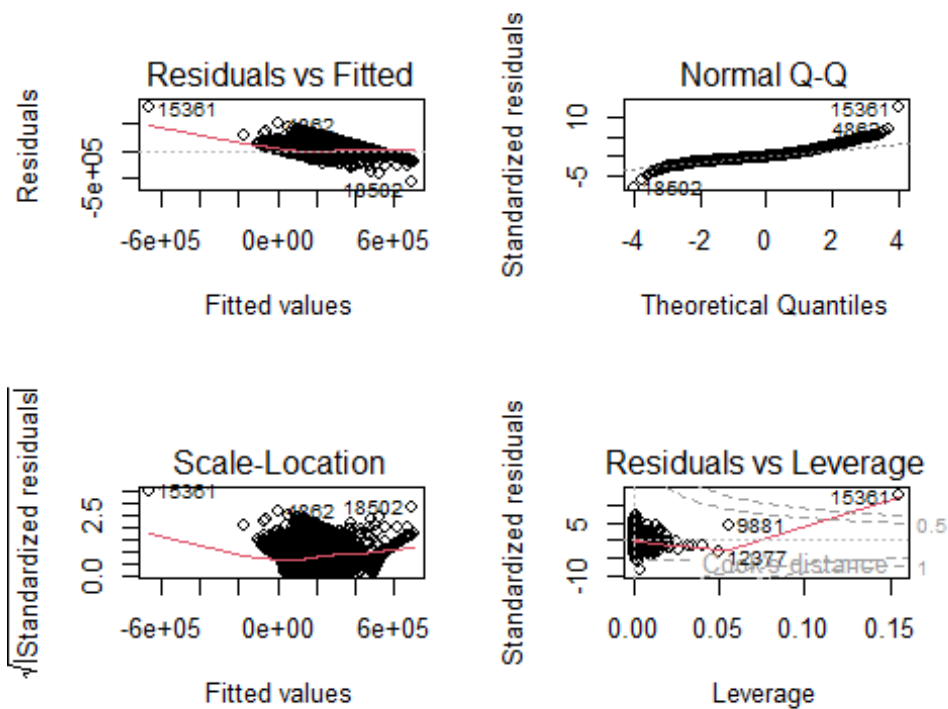
## [1] 4597020605

print(rmse1 <- sqrt(mse1)) # The RMSE.

## [1] 67801.33

par(mfrow = c(2, 2))
plot(lm1)

```



kNN Regression

```
library(caret)

## Warning: package 'caret' was built under R version 4.2.3

## Loading required package: lattice

# Model fitting.
fit <- knnreg(train[,3:9], train[,2], k = 5)

# Evaluation.
pred2 <- predict(fit, test[,3:9])
print(cor_knn1 <- cor(pred2, test$median_house_value)) # Correlation.

## [1] -0.2835157

print(mse_knn1 <- mean((pred2 - test$median_house_value)^2)) # The MSE.

## [1] 56220849031

print(rmse2 <- sqrt(mse_knn1)) # The RMSE.

## [1] 237109.4

# Data scaling.
train_scaled <- train[,3:9]
means <- sapply(train_scaled, mean)
stdvs <- sapply(train_scaled, sd)
```

```

train_scaled <- scale(train_scaled, center = means, scale = stdvs)
test_scaled <- scale(test[,3:9], center = means, scale = stdvs)

# knn for the scaled data.
fit <- knnreg(train_scaled, train$median_house_value, k = 7)
pred3 <- predict(fit, test_scaled)
print(cor_knn2 <- cor(pred3, test$median_house_value))

## [1] 0.9947504

print(mse_knn2 <- mean((pred3 - test$median_house_value)^2))

## [1] 143304622

print(rmse3 <- sqrt(mse_knn2)) # The RMSE.

## [1] 11970.99

# Finding the best k algorithm.

cor_k <- rep(0, 25)
mse_k <- rep(0, 25)

num <- 1
for (k in seq(1, 43, 2))
{
  fit_k <- knnreg(train_scaled, train$median_house_value, k = k)
  pred_k <- predict(fit_k, test_scaled)
  cor_k[num] <- cor(pred_k, test$median_house_value) # Correlation value.
  mse_k[num] <- mean((pred_k - test$median_house_value)^2) # The MSE.
  print(paste("k = ", k, cor_k[num], mse_k[num]))
  num <- num + 1
}

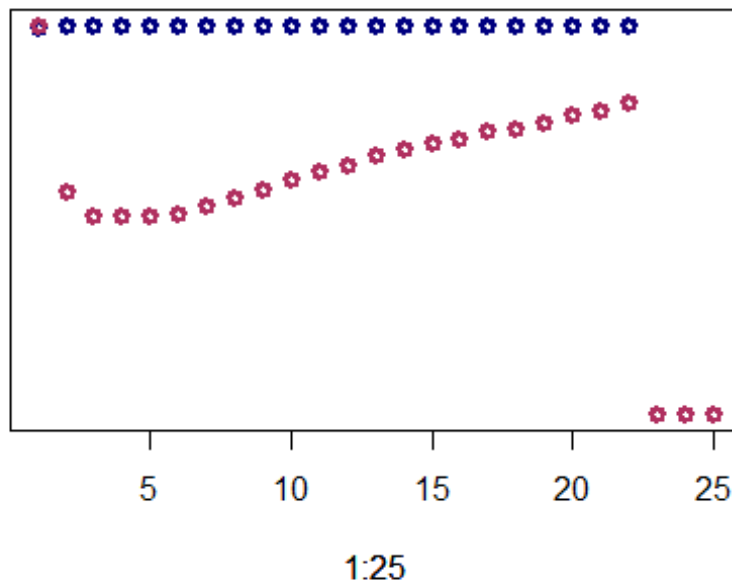
## [1] "k = 1 0.989279387434334 282376990.03034"
## [1] "k = 3 0.993897558621111 162107770.556398"
## [1] "k = 5 0.994639521093658 144492354.970546"
## [1] "k = 7 0.994750383640009 143304621.529022"
## [1] "k = 9 0.994778005188362 143773832.343323"
## [1] "k = 11 0.994747082359241 145897116.989246"
## [1] "k = 13 0.994598027558272 151317348.69291"
## [1] "k = 15 0.994436361673998 156506027.781358"
## [1] "k = 17 0.994272519842331 162819376.379659"
## [1] "k = 19 0.994052564769095 170015279.548431"
## [1] "k = 21 0.993860263921137 176852526.423217"
## [1] "k = 23 0.993758472978828 181382385.607802"
## [1] "k = 25 0.99355383593146 188172512.248087"
## [1] "k = 27 0.993460491962776 192305629.276772"
## [1] "k = 29 0.993347069622814 196723299.419251"
## [1] "k = 31 0.993287721370641 199973395.641524"
## [1] "k = 33 0.993151928890091 205115339.198631"

```

```
## [1] "k = 35 0.993116345580429 207652602.101231"
## [1] "k = 37 0.992994140712848 211966216.934172"
## [1] "k = 39 0.992861386146755 217211630.00099"
## [1] "k = 41 0.992814959484625 220293272.57596"
## [1] "k = 43 0.992668002402656 225787838.949563"

# Plotting the kNN regression model.
plot(1:25, cor_k, lwd = 3, col = 'navy', ylab = "", yaxt = 'n')
par(new = TRUE)
plot(1:25, mse_k, lwd = 3, col = 'maroon', labels = FALSE,
     ylab = "", yaxt = 'n')

## Warning in plot.window(...): "labels" is not a graphical parameter
## Warning in plot.xy(xy, type, ...): "labels" is not a graphical parameter
## Warning in box(...): "labels" is not a graphical parameter
## Warning in title(...): "labels" is not a graphical parameter
```



```
# Lowest and highest values.
which.min(mse_k) # The lowest.

## [1] 23

which.max(cor_k) # The highest.

## [1] 5
```

As seen from the algorithm and diagram above, the lowest MSE came from “ $k = 23$ ” and the highest correlation coefficient came from “ $k = 5$.”

Comparison & Analysis (Parts C & D)

Given the regression techniques which I implemented, the correlation coefficients and mean squared errors were varied. Logically speaking, this is expected in the sense that which kNN generally had higher correlation coefficients of 99% whereas the linear regression technique only had one of roughly 80%, meaning that for the data set, kNN is superior in this case. Also, kNN had a lower mean squared error than linear regression.

Analyzing the techniques, there are clearly varying methods of retrieving the results as seen from the techniques above. Additionally, kNN has much more flexibility in terms of the methodologies the technique can process as the possibilities are theoretically endless. This is due to the fact that it does not that there is a response variable, per se. Conversely, linear regression is much simpler to utilize as there is a general approach to how the model will fit the data, given that there is not only a target, but also the predictors which contribute towards the target variable.